Munzur Atak

AI & Software Engineer | LLM Systems, RAG & Evaluation | Python, TypeScript, Azure & Microsoft 365

+31 6 39009964 | munzuratak@outlook.com | linkedin.com/in/munzur-atak-284b62225 | github.com/MunzurAtak

Rotterdam, Netherlands | open to hybrid roles in the Randstad

Final-year BSc Artificial Intelligence student at VU Amsterdam (graduating December 2026) and working student Data & AI at Morgens Conclusion, a consultancy with 100+ consultants. I build LLM systems end to end: retrieval pipelines, LoRA fine-tuning and the evaluation that proves they work. I also develop internal AI applications, Copilot Studio agents and Power BI reporting used up to board level. Thesis on stance preservation in multi-agent LLM debates graded 8.5. **Looking for a full-time AI engineering role from January 2027.**

TECHNICAL SKILLS

AI & LLM Engineering: RAG pipelines, vector search (FAISS), LLM fine-tuning (LoRA, QLoRA), LLM evaluation (LLM-as-a-judge, annotator agreement, Cohen's kappa), prompt & agent design, local model inference (Ollama), NLP, text classification

Machine Learning: PyTorch, Hugging Face (Transformers, PEFT), scikit-learn, model training & evaluation, feature engineering

Languages: Python, SQL, TypeScript, JavaScript

Data & Business Intelligence: Power BI (data modelling, reporting, dashboard design), SQL, pandas, NumPy

Cloud & Microsoft platform: Microsoft Copilot Studio, Power Automate, SharePoint, Microsoft 365, Microsoft Azure and Azure AI Foundry (foundational), Supabase, Cloudflare

Software Engineering: FastAPI, React, Vite, REST APIs, Docker, Git, GitHub Actions (CI/CD), Pytest

EXPERIENCE

Working Student Data & AI - Morgens Conclusion Jun 2026 - present

Management and digital consultancy, 100+ consultants, part of the Conclusion group

- **Internal AI knowledge app:** building an application that helps 100+ consultants find the methods and information they need for client work. Built with TypeScript, React and Vite, running in preview on Cloudflare and Supabase, and being prepared for production within Morgens' Microsoft environment.
- **Copilot Studio agent:** built and rolled out an AI agent that automates drafting project plans and emails the result to the user.
- **AI team:** part of the team responsible for AI tooling at Morgens, co-responsible for the internal Claude pilot, and supporting colleagues with AI and data questions in client engagements.
- **Healthcare analytics (WeConnect):** co-developed Power BI dashboards for four programmes within WeConnect, based on data from hospitals in the Rijnmond region. Board members and project leads use them to decide on software adoption in their organisations.
- **Business intelligence:** built three business-critical Power BI dashboards with finance and sales insights for consultants and the management team.

Business Operations & Internal Automation - 010 Op Niveau Apr 2025 - Sep 2026

Tutoring organisation, ~30 staff, Rotterdam

- Automated monthly payroll by building, with Claude, a tool that generates SEPA batch files (pain.001.001.03) directly from exported time registrations, replacing manual entry per employee.
- Built and maintained the organisation's website, including a Typeform-based CITO level test in the student intake.
- Processed student and family balance administration from CSV exports, with anomaly checks before invoicing.

Administrative Officer & Director Support - ProTeam BV May 2023 - May 2024

Security company, Rotterdam

- Supported the director with reporting, documentation and operational logistics.

RESEARCH

BSc thesis: Comparing LLM Methods for Stance Preservation in Adversarial Debates 2026

Vrije Universiteit Amsterdam | Supervisor: Filip Ilievski | Grade: 8.5 |
github.com/MunzurAtak/stance-preservation-llm-debates

- Compared prompting, stance-anchored RAG and LoRA fine-tuning on their ability to hold an assigned stance across 90 controlled five-round adversarial LLM debates, totalling 900 generated turns.
- Built the full pipeline in Python: local multi-model inference via Ollama (Qwen2.5-7B, Llama 3.1 8B, Gemma 3 12B), a FAISS retrieval index with sentence-transformers embeddings, and a QLoRA adapter trained on 1,500 examples.
- Designed the LLM-as-a-judge evaluation and validated the instrument itself against one human and two external LLM annotators using Cohen's kappa, reaching 93-100% polarity agreement.
- Concluded that external grounding via RAG preserves the assigned stance most reliably, and that most failures are driven by the normative direction of the stance rather than by the conditioning method.

PROJECTS

ContextAI - Retrieval-Augmented Generation Service 2026

github.com/MunzurAtak/contextai-rag-service

- RAG service where users upload a PDF and ask questions about it. Two-stage retrieval pipeline with FAISS bi-encoder search followed by cross-encoder reranking, answer generation via Llama 3.1 on the Groq API, and confidence scoring based on the reranker.
- Built with FastAPI and containerised with Docker Compose, with structured logging and per-stage latency metrics (embedding, search, reranking, generation) on every request.

ML Inference API - Dockerised FastAPI Service 2026

github.com/MunzurAtak/ml-inference-api

- Production-style machine learning inference REST API built with scikit-learn and FastAPI, with Pytest test coverage and end-to-end Docker packaging.

EDUCATION

BSc Artificial Intelligence - Vrije Universiteit Amsterdam expected Dec 2026
VWO Gymnasium (pre-university) - Marnix Gymnasium 2021

Languages: Dutch (native), Turkish (native), English (fluent)